

Performing *embarrassingly* parallel data analysis in Python using **ipyparallel** and **pandas**

Marcos Oliveira

Why?

- An embarrassingly parallel problem

A problem that may be solved in parallel by easily splitting it into separate problems.

- In data analysis, quite often:
 - Same procedure for different parts of the data:

Example: Last.fm data set

	user_id	artist	album	song	time_stamp	datetime
0	31435741	2	4	4	1385212958	2013-11-23 13:22:38
1	31435741	2	4	4	1385212642	2013-11-23 13:17:22
2	31435741	2	4	4	1385212325	2013-11-23 13:12:05
3	31435741	2	4	4	1385209508	2013-11-23 12:25:08
4	31435741	2	4	4	1385209191	2013-11-23 12:19:51

Example: a Twitter data set

	user_id	datetime	link
0	140915906	2010-10-01 16:48:25	http://kingo.to/aY9
1	32253363	2010-09-26 18:00:47	http://www.aquapropertiesinc.com
2	170948166	2010-09-24 05:40:14	http://twitpic.com/2r8l0d
3	163171731	2010-09-23 16:21:54	http://lc4.in/6GeZ
4	105539322	2010-09-23 19:49:51	http://migre.me/1iQxq

Why?

- Embarrassingly parallel data analysis

Available processing units

1	[0.0%	15	[3.7%	29	[0.0%	43	[0.0%
2	[0.0%	16	[0.6%	30	[0.0%	44	[0.0%
3	[0.0%	17	[0.0%	31	[0.0%	45	[0.0%
4	[0.0%	18	[0.6%	32	[0.0%	46	[0.0%
5	[0.0%	19	[1.2%	33	[23.8%	47	[0.0%
6	[0.0%	20	[1.8%	34	[0.0%	48	[0.0%
7	[0.0%	21	[0.0%	35	[0.0%	49	[0.0%
8	[0.0%	22	[0.0%	36	[0.6%	50	[0.0%
9	[9.1%	23	[0.0%	37	[0.0%	51	[0.0%
10	[0.6%	24	[2.4%	38	[0.0%	52	[0.6%
11	[0.0%	25	[1.2%	39	[0.6%	53	[0.0%
12	[0.0%	26	[12.3%	40	[0.0%	54	[0.6%
13	[4.2%	27	[0.6%	41	[0.0%	55	[0.6%
14	[5.5%	28	[0.0%	42	[0.6%	56	[0.0%

Usage - what we want:

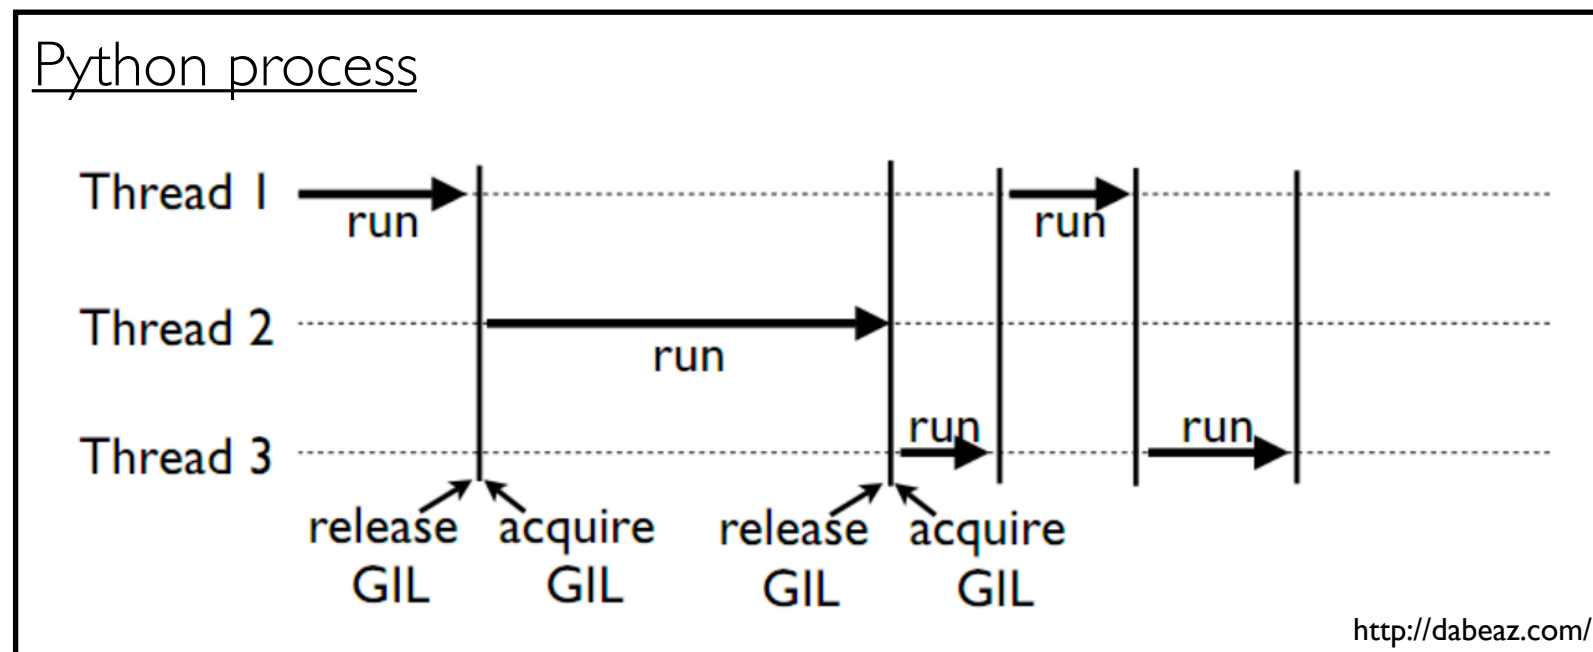
1	[92.5%	15	[75.2%	29	[39.0%	43	[88.1%
2	[49.4%	16	[66.7%	30	[85.6%	44	[92.5%
3	[85.6%	17	[61.1%	31	[38.8%	45	[87.5%
4	[75.6%	18	[70.2%	32	[45.9%	46	[91.9%
5	[52.9%	19	[70.6%	33	[67.1%	47	[63.8%
6	[57.5%	20	[62.1%	34	[72.5%	48	[86.4%
7	[71.9%	21	[66.0%	35	[50.0%	49	[92.5%
8	[93.7%	22	[66.2%	36	[55.6%	50	[92.5%
9	[60.1%	23	[66.7%	37	[41.1%	51	[85.8%
10	[58.8%	24	[75.6%	38	[66.2%	52	[70.2%
11	[67.1%	25	[87.6%	39	[45.6%	53	[55.7%
12	[73.1%	26	[79.1%	40	[48.8%	54	[69.9%
13	[48.4%	27	[70.0%	41	[91.7%	55	[84.4%
14	[96.9%	28	[65.2%	42	[16.9%	56	[87.6%

Usage - what we sometimes have:

1	[6.4%	15	[0.0%	29	[0.0%	43	[0.0%
2	[9.7%	16	[100.0%	30	[0.0%	44	[0.0%
3	[4.8%	17	[0.0%	31	[0.0%	45	[0.0%
4	[1.2%	18	[0.6%	32	[0.0%	46	[0.0%
5	[0.6%	19	[0.0%	33	[0.0%	47	[21.8%
6	[0.0%	20	[0.6%	34	[0.0%	48	[0.0%
7	[1.2%	21	[0.0%	35	[0.0%	49	[0.0%
8	[4.2%	22	[0.0%	36	[3.7%	50	[0.0%
9	[8.5%	23	[0.0%	37	[0.0%	51	[0.0%
10	[1.8%	24	[0.0%	38	[0.0%	52	[0.0%
11	[1.2%	25	[0.0%	39	[5.5%	53	[0.0%
12	[0.0%	26	[0.0%	40	[0.6%	54	[0.0%
13	[5.5%	27	[0.0%	41	[0.0%	55	[0.0%
14	[0.0%	28	[0.6%	42	[0.6%	56	[0.0%

Python

- Python allows multi-thread *programming*... `import thread`
- but GIL — the global interpreter lock



- GIL prevents multi-thread *execution*.
- “*Solution*”: instead of multi-threads, multi-processes.
 - Reminder: CPU bounded vs. I/O bounded

The good the bad the ugly

- Multiprocessing native in Python

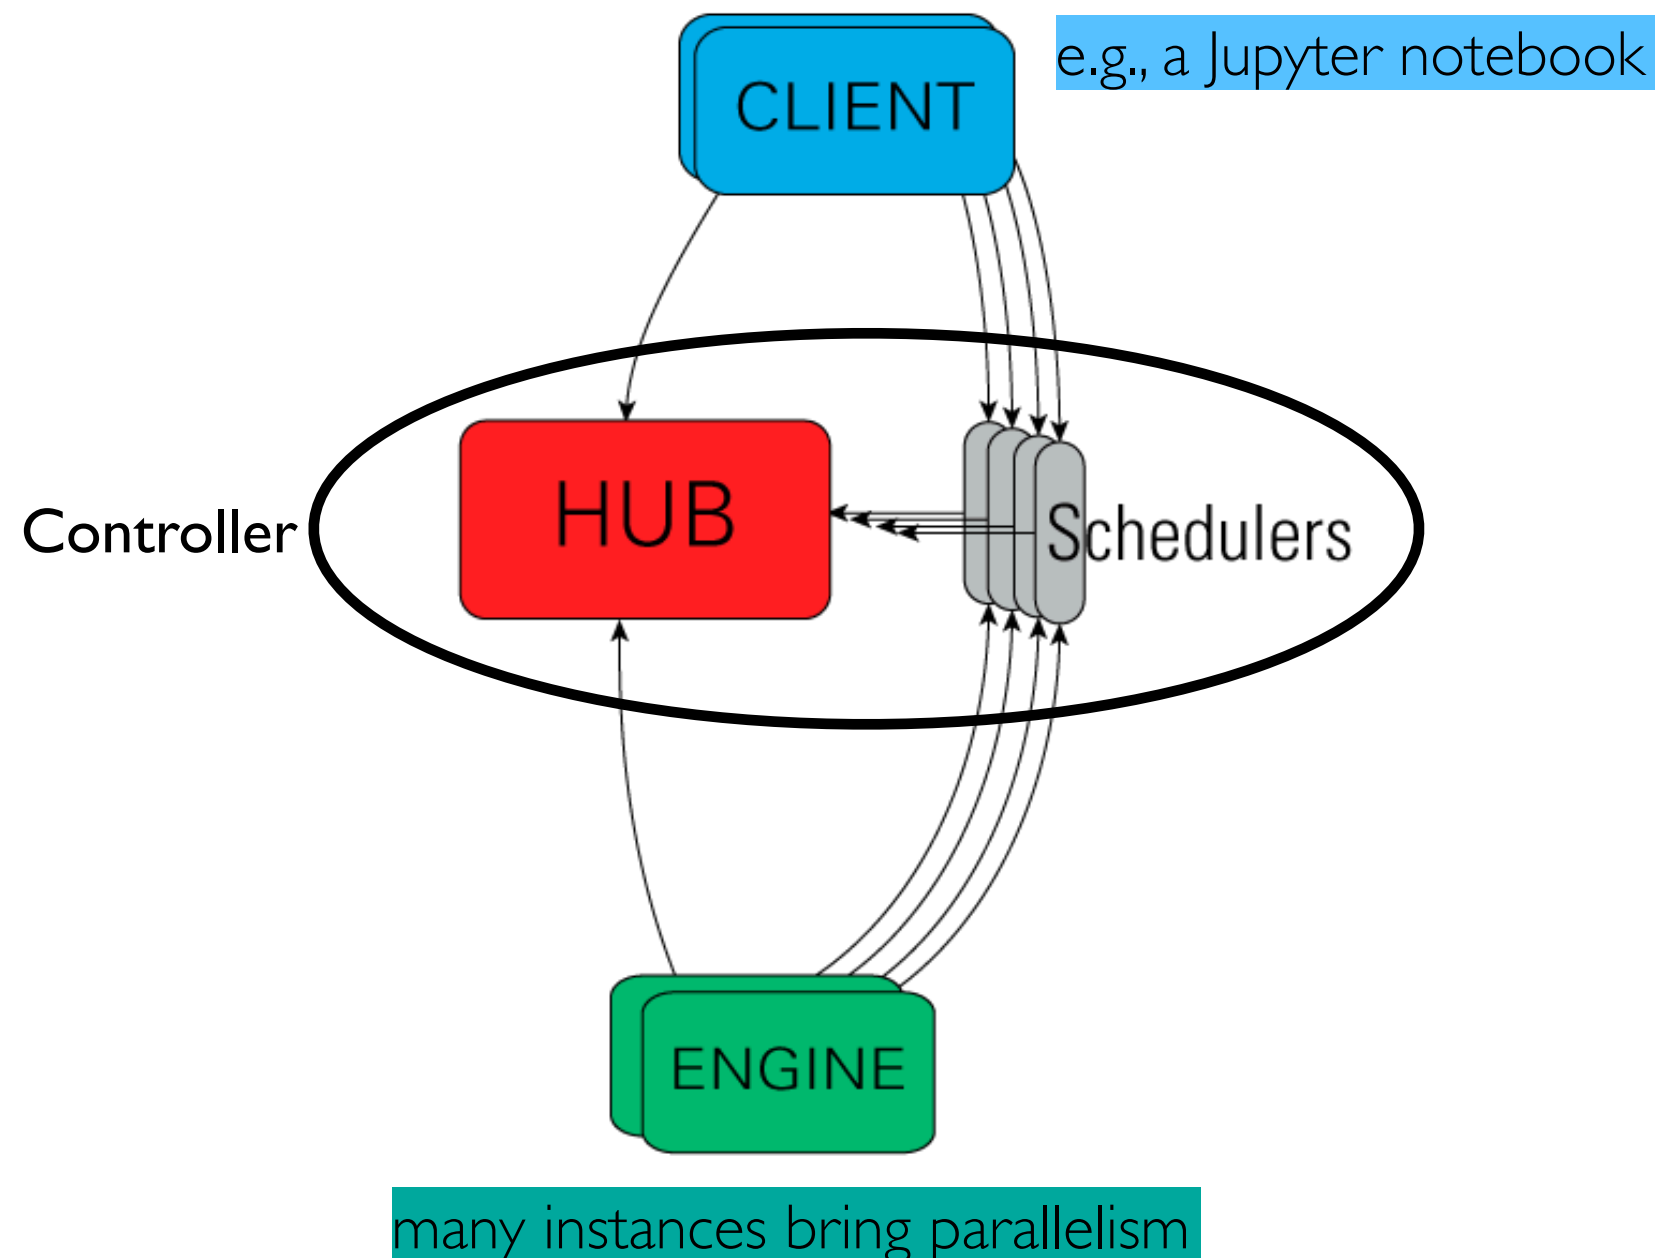
```
from multiprocessing import Pool
```

```
from multiprocessing import Process
```

- IPyparallel
 - Jupyter notebook ~ quick-dirty coding, preliminary analyses
 - With pandas, we have parallel quick data analysis in few code lines.
 - Interactive

IPyparallel

- In a nutshell:



IPyparallel

- Installing via pip:

```
pip install ipyparallel
```

- Creating controller + 4 engines:

```
ipcluster start -n 4
```

- By the way, ipcluster is a shortcut that creates both controller and engine in the localhost.
 - You can have more control of configuration using profiles:
<https://ipyparallel.readthedocs.io/en/latest/process.html>

IPyparallel

- Coding time